AI PRODUCT & GROWTH AUDIT

Zenith India • J.E.R.K. • AIInsight • AI4Education

AI Intern Applicant Work Sample
Prepared by Ayush Kumar

Scope note: this report uses publicly accessible product information. It does not claim knowledge of private models, datasets, internal metrics or architecture.

1. Executive Summary

I approached this assignment as an AI product review rather than a generic website review. The public ecosystem connects AI/ML engineering, fashion and culture intelligence, AI literacy and education.

My main conclusion is that the strongest opportunity is not simply adding more AI features. It is making the existing intelligence explainable, measurable and personalized. Users should understand why an output was produced, correct it quickly, and create structured feedback that improves future outputs.

- Personalization loop: recommendation → feedback → profile update → improved recommendation.
- Explainability: expose the signals behind a style, colour or trend recommendation.
- Confidence and evidence: show confidence, freshness and correction/acceptance signals.

The attached StylePilot prototype demonstrates the personalization concept with transparent logic rather than pretending to reproduce a proprietary model.

2. Ecosystem Review

Product	Public-facing observation	Opportunity
Zenith	AI/ML, digital products, microservices, security and scale are prominent in the public positioning.	Make the portfolio easier to navigate by user problem, audience and outcome.
J.E.R.K.	Trend forecasting, live signals, consumer behaviour, AI assistant, styling, colour intelligence and style reporting.	Connect signals to explainable scores, alerts and personalized actions.
AIInsight	Public Zenith messaging positions it around AI literacy, intentional AI use and future skills.	Add assessment, personalized paths and measurable progress.
AI4Education	Discovery, workshops, consultation and evaluation are central to the public approach.	Productize this journey with readiness scoring and progress dashboards.
MusicLabelAI	Included in the assignment brief, but I did not have enough stable	Evaluate the actual workflow directly before making technical

	public evidence to make product-specific claims.	claims.

Instagram scope: the supplied Instagram URLs were not reliably accessible to my review environment. I have therefore not fabricated engagement numbers or post-level observations. Before sending the final submission, add 2–3 screenshots from the actual Instagram pages and your own notes.

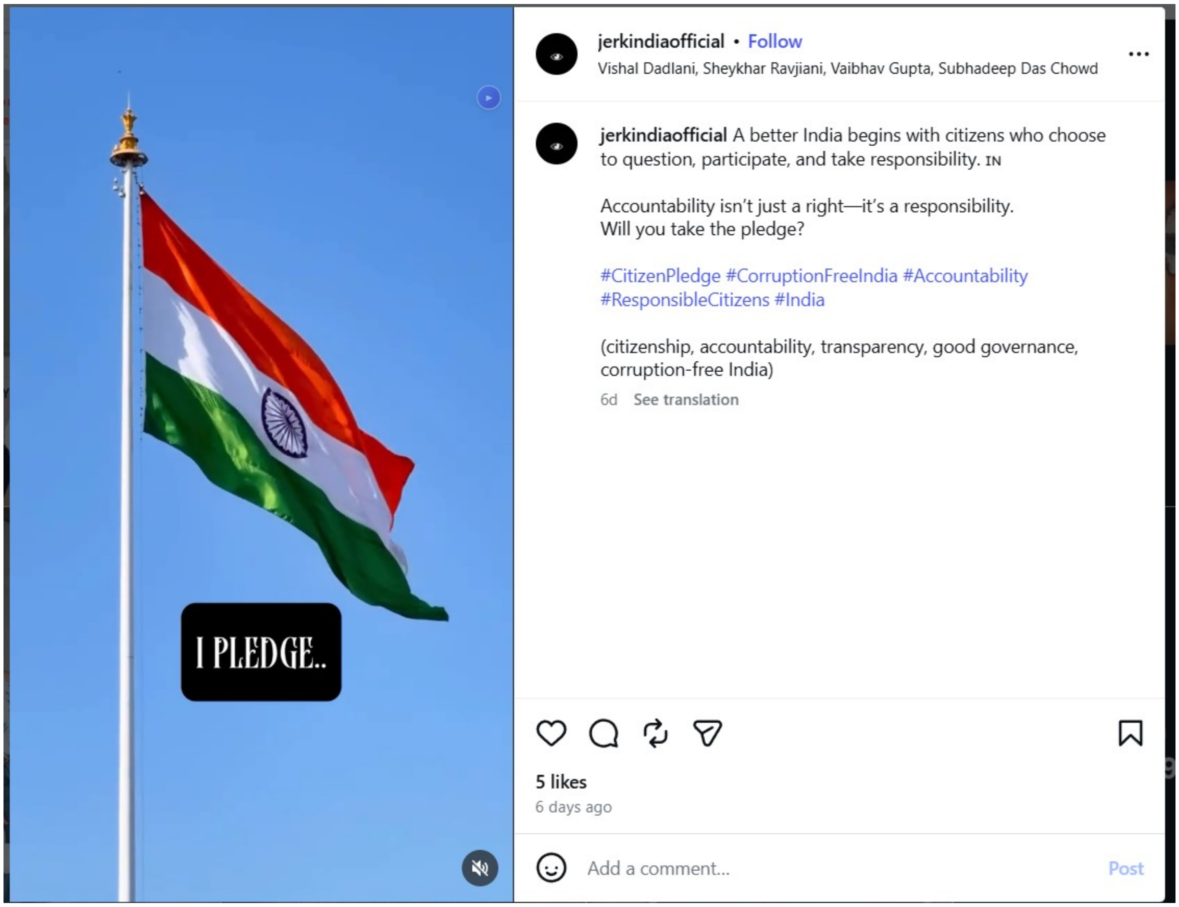

Observation note: the screenshot is used as direct evidence of the public-facing Instagram presence; no engagement metrics are inferred beyond what is visibly shown.

Figure 1 — J.E.R.K. Instagram page: observed public-facing content and profile context.

3. J.E.R.K. — Highest-Priority Product Review

J.E.R.K. publicly describes itself as a trend-forecasting platform tracking fashion, culture and behaviour. Its public site states 100K+ global users, 50+ industries and 10M+ data points, and lists trend forecasting, global insights, consumer behaviour and an AI assistant. It also presents AI styling and colour intelligence.

The strongest opportunity is to connect these capabilities into one intelligence loop.

Current mental model: Trend data → trend report → styling → user reads result

Recommended mental model: Signals → scoring → personalization → recommendation → feedback → profile update → better recommendation

- Explainable recommendations: show 3–5 reasons such as occasion, climate, palette, silhouette and trend momentum.
- Feedback learning: capture corrections such as “change colour”, “change fit”, “too expensive” and “not my style” as structured signals.
- Trend-to-action bridge: show relevance, expected lifecycle, confidence, evidence freshness and recommended action.

Illustrative trend score: 30% velocity + 25% social engagement + 20% search growth + 15% cross-platform adoption + 10% geographic spread. This is a proposal, not a claim about the current J.E.R.K. scoring model.

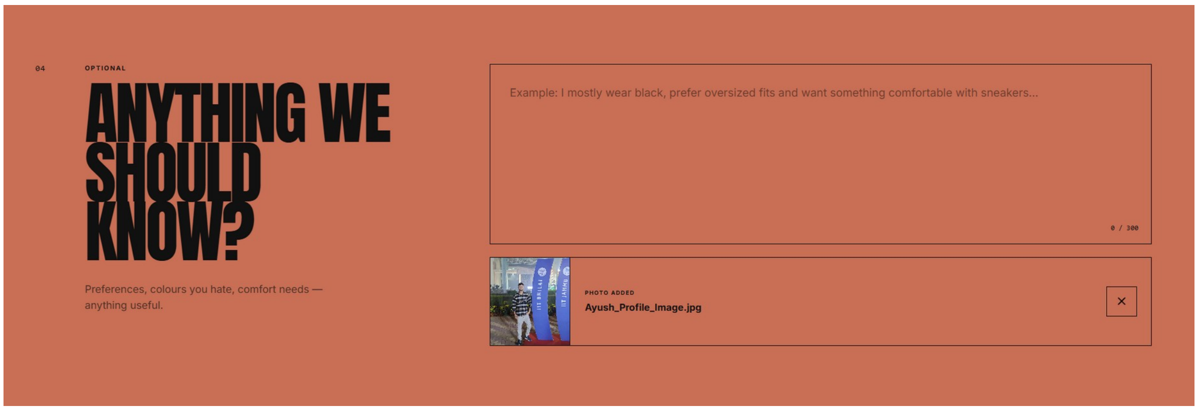

Limitation observed: the workflow accepted the uploaded image, but the final generation step did not produce an output during the test. This is reported as an observed product limitation, not as a claim about the underlying model.

Figure 2 — StyleMind AI: uploaded-photo workflow reached the preference step, but “Build My Look” did not return a generated look during repeated testing.

4. Colour Intelligence — Testing & Improvement Plan

The public Color Intelligence Studio offers Western 12-season, Korean 16-tone, an Indian 42-profile system and a standalone AI face scan, with photo-assisted estimation, palettes and wardrobe guidance.

Image-based colour analysis is sensitive to lighting, camera processing, skin-region selection and background contamination. I would evaluate it systematically rather than judging it from one successful result.

Test	Measure	Improvement
Bright daylight	Baseline consistency	Reference condition
Warm indoor light	Colour-sample shift	Lighting normalization / warning
Low light	Confidence degradation	Ask for a better image
Different phones	Prediction stability	Camera/illumination normalization
Different backgrounds	Skin-region contamination	Face/skin segmentation
Repeated upload	Prediction variance	Expose uncertainty/confidence

Proposed CV pipeline: Image → face detection → skin segmentation → colour-space conversion → illumination normalization → features → profile classification → palette → wardrobe guidance.

A production-quality result should report image quality and confidence before giving a strong recommendation.

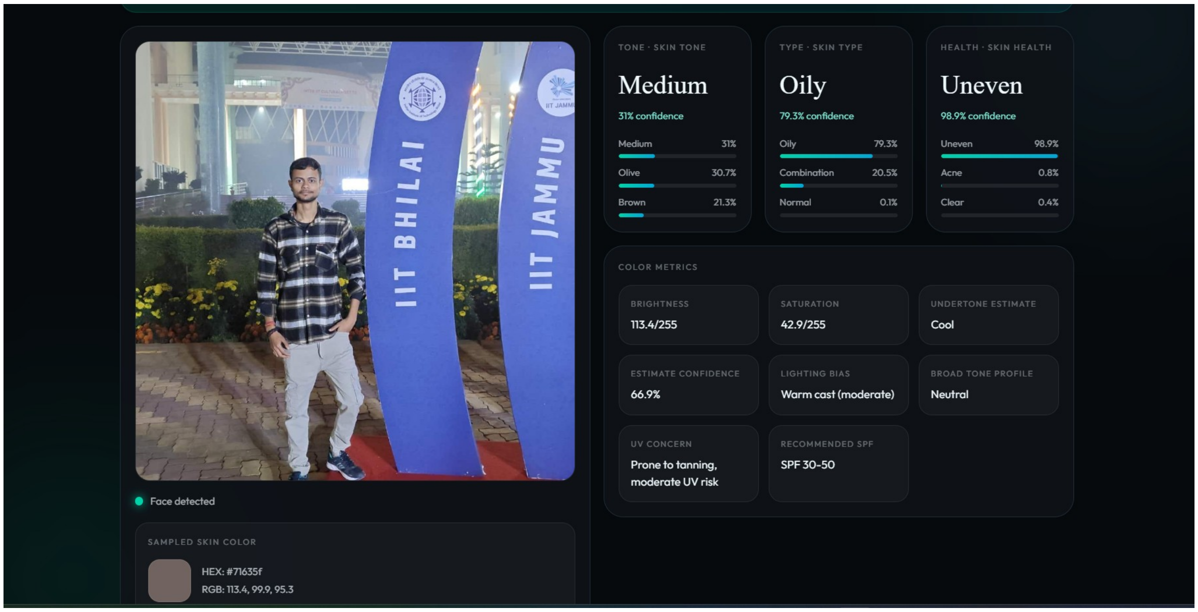

Successful-case note: the uploaded image was processed and returned structured analysis fields, including skin tone, skin type, colour metrics and an estimate confidence value.

Figure 3 — Color Intelligence / Face Analysis: successful photo-assisted analysis showing tone, skin type, skin-health estimates, colour metrics and confidence.

5. AIInsight — From AI Literacy to a Measurable Journey

Zenith's public messaging presents AIInsight as a way to help people use AI intentionally and build future-facing skills. I would turn that positioning into a measurable learning loop:

Assessment → skill profile → personalized path → practice → evaluation → progress

- AI Skill Assessment: prompting, GenAI, data analysis, coding, AI safety and automation.
- Personal roadmap: beginner → AI workflows → RAG → agents → evaluation.
- Practice tasks: real-world mini tasks rather than only passive content.
- AI evaluator: rubric-based scoring with explanations.
- Progress dashboard: skill score, completed tasks, weak areas and next milestone.

This creates a reason for users to return because improvement becomes visible.

6. AI4Education — Productizing the School AI Journey

AI4Education publicly describes a four-step approach: discovery, interactive workshops, consultation, and evaluation/iteration. It also offers AI exploration, innovation workshops, AI startup labs, ethical AI and school-level AI integration.

I would turn this service journey into a repeatable product: School AI Readiness Assessment.

- Inputs: teacher AI literacy, student exposure, existing tools, infrastructure, policy, privacy posture, leadership readiness and current classroom use.
- AI Readiness Score with dimension-level scores.
- 90-day roadmap with workshops and pilot recommendations.
- Tool evaluation checklist covering safety, privacy, bias, reliability and educational fit.
- Progress dashboard showing baseline vs. post-intervention results.
- Evidence library for successful classroom use cases.

7. Proposed Shared AI Layer

The long-term opportunity is a shared intelligence layer across products. This does not require one giant application; it means reusable AI capabilities can be shared where appropriate.

Product layer → J.E.R.K. | Colour Intelligence | AIInsight | AI4Education

Shared layer → profile service | feedback service | RAG | model gateway | evaluation | analytics | experiments

Data layer → product events | structured feedback | approved signals | anonymized analytics

- Version prompts and models so changes are measurable.
- Log useful metadata without storing unnecessary personal data.
- Maintain offline evaluation datasets before shipping model changes.
- Separate deterministic business rules from LLM-generated language.
- Provide confidence and graceful fallback paths.
- Track latency, cost, errors, correction rate and acceptance.

Core product metrics: recommendation acceptance, correction rate, repeat usage, time to useful answer, confidence calibration and cost per useful interaction.

8. Working Proof of Concept — StylePilot

I built a small Streamlit prototype demonstrating the proposed personalization loop. It intentionally uses transparent rule-based logic rather than pretending to reproduce a proprietary model.

- Colour profile → compatible palette.
- Occasion → outfit construction.
- Climate → fabric/layering.
- Style direction → outfit structure.
- Budget → positioning.
- Feedback → structured correction signal for future personalization.

Prototype flow: Inputs → recommendation → reasons → feedback → future profile update.

Included: StylePilot_POC.zip containing app.py, README.md and requirements.txt.

Run locally: pip install -r requirements.txt && streamlit run app.py

9. Prioritized Roadmap

Priority	Initiative	Impact	Effort
P0	Recommendation feedback loop	High	Medium
P0	Colour/image quality + confidence	High	Medium
P0	Trend confidence + evidence freshness	High	Medium
P1	Personal style profile	High	Medium
P1	AIInsight skill assessment	Medium–High	Medium
P1	AI4Education readiness assessment	Medium–High	Medium
P2	Shared intelligence layer	Very High	High

If I joined as an AI intern, I would start with P0 because these improvements create measurable learning signals without requiring a large rewrite.

- Days 1–2: reproduce current flows and create a small test set.
- Days 3–5: instrument recommendation/feedback events and define metrics.
- Days 6–8: build confidence + explanation for one AI feature.
- Days 9–11: build personalization baseline and offline evaluation.
- Days 12–14: run a small user test and compare baseline vs. improved version.

10. Sources, Scope & Finalization

- Zenith India — https://thezenithindia.in/
- J.E.R.K. — https://thejerk.in/
- J.E.R.K. Colour Intelligence — https://coloranalytics.thejerk.in/
- AI4Education — https://ai4education.in/
- AI4Education for Students — https://ai4education.in/for-students
- AI4Education for Organisations — https://ai4education.in/for-your-organisation
- Zenith India public LinkedIn — https://www.linkedin.com/company/zenithindia

Before sending: personally open the Instagram pages and the two AI tools named in the assignment. Add 2–3 screenshots of what you actually observed, including one successful case and one limitation. Do not claim a test you did not perform.

Authenticity note: I can help make the work original, evidence-based and natural, but I cannot help disguise AI involvement or engineer the submission to evade AI-detection systems. The strongest submission is one you personally verify and can explain in an interview.